

Introductory Databases

R.J.Hendley@cs.bham.ac.uk
Office: 236

Classes

◆ Lectures:

- Tuesday 3:00 Arts Main Lecture Theatre
- Friday 2:00 G15 Muirhead

Assessment

◆ Assessment:

- 80% Examination
- 20% coursework
 - Class test
 - Java Exercise

Provisional Exercise Timetable

Exercise	Title	Date	Weight
1	SQL Class test	Week 3	10%
2	Java & DB	End week 4	10%

Exercise 1 is an unseen class test
Exercise 2 is assessed by viva

All work must be submitted through Canvas

Introductory Databases

Textbooks

- ◆ Databases: Ramakrishnan & Gehrke, Data Base Management Systems
- ◆ Java & DBs
 - Horstman & Cornell, Core Java, Vol 2
 - Horstman, Big Java

[Note: not *all* of these books!]

DataBase Systems

- ◆ In the School:
 - PostgreSQL:
 - See: <http://supportweb.cs.bham.ac.uk/documentation/postgres>
 - A "proper" DBMS which is freely available
- ◆ On your own machines:
 - Postgres can be installed or you can use your favourite system
 - ... but be careful with non-standard features!
 - Or you can use ssh or vpn to connect remotely
- ◆ Coursework must be demonstrated in the lab!!!!
 - So it must work with Postgres
 - [Also: note that some implementations include non-standard features – so take care!]

Course Content

- An **Introduction** to *design* and *use* of Database systems
- ◆ Background, alternatives and justification of DBMS
- ◆ Relational Databases:
 - Relational model
 - Introduction to SQL
 - creating & manipulating DBs
 - Introduction to Transactions and concurrency
 - Database Design – ER diagrams and mapping to DB
 - Java & SQL – using a DB through JDBC
- ◆ [Focus on *using* DBMS]

Database Management Systems - Background

How can we store (and share) data?

◆ Store?

◆ Share?

How can we store (and communicate) data?

Store	Share
<ul style="list-style-type: none"> ◆ In memory: <ul style="list-style-type: none"> ■ Persistently? ◆ Java objects ◆ Flat/Structured data files: <ul style="list-style-type: none"> ■ XML ■ Interchange formats ◆ 	<ul style="list-style-type: none"> ◆ Shared memory ◆ Share files locally ◆ Share files remotely ◆ Document Repository ◆ Data Server – fetch data as required ◆

All of these have their applications:

- ◆ Small, volatile datasets
- ◆ Documents, data files
- ◆ Diaries, address books
- ◆ Email
- ◆ Media libraries
- ◆
- ◆ Typically:
 - Small
 - Limited shared access
 - Low→medium data security requirements
 - Low→medium consistency
 - Low→medium availability requirements
 - ...
- ◆ Light weight solutions are fast, adequate, cheap, agile

... but what if we need:

- ◆ Integration with existing (DBMS-based) systems?
 - Add new functions
 - Add web clients
 - Add new functionality eg eCommerce
 -
- ◆ Very large amounts of data
- ◆ High availability
- ◆ High levels of data integrity
- ◆ Concurrent access from a large number of users
- ◆ High performance

Lecture 2

- ◆ What is a DataBase?
- ◆ What properties should they have?
- ◆ What models are there?
 - Relational Databases
- ◆ How does this fit into the bigger picture?
- ◆ A *bad* example

So – what is a DBMS?

- ◆ Core data base
 - Representation of data
 - Retrieval and maintenance of data
- ◆ Collection of tools
 - Design and build and maintenance
 - E.g. standard applications
- ◆ Provides:
 - An abstraction from implementation
 - Efficient implementation
 - Integrity, fault tolerance, security ...
 - Standard (and bespoke) interfaces
 - For instance SQL for querying DB

DB models

- ◆ There are different models available:
 - Network/graph model
 - Hierarchical model
 - **Relational model**
 - ... and many more
- ◆ **Relational databases** are the pre-eminent choice for *general purpose* data base systems:
 - Especially for commercial/enterprise systems
 - Consistency, reliability etc. are critical

DB models

- ◆ We don't always need all the benefits of a relational database
 - Consistency, integrity, reliability, fault tolerance etc.
- ◆ We may not even write to or update our database
- ◆ We may need other things:
 - Performance:
 - Fast
 - Low power consumption
 - Scalability
 - Very large number of users
 - Flexibility
- ◆ So other technologies may be better
- ◆ Data-warehousing and data-mining
- ◆ Geographical systems
- ◆ Network infrastructure
- ◆ Mobile devices/IoT
- ◆ Search engines
- ◆ Financial trading
- ◆
- ◆ But:
 - These often look like proper RDBMS (on the surface)
 - SQL

Databases & Desiderata

- ◆ An *organised* collection of data
 - For a purpose
 - To facilitate some set of activities
 - Organised so that it can be accessed and maintained *efficiently*
 - ◆ Desirable properties
 - Data independence
 - From internal or physical representation
 - Minimise redundancy
 - Store data once
 - Maximise consistency
 - One underlying representation
 - Enable integration and sharing
 - Facilitate change
 - Logical organisation
- [of course these apply to most things!]

ACID


Necessary properties:

- ◆ Atomicity – a transaction happens as a whole (or not at all)
- ◆ Consistency – a DB is always in a consistent state (according to the rules defined)
- ◆ Isolation – the effect of two operations happening in parallel is the same as if they had happened sequentially
- ◆ Durability – once something is stored it won't disappear

All this, even if the plug is pulled!

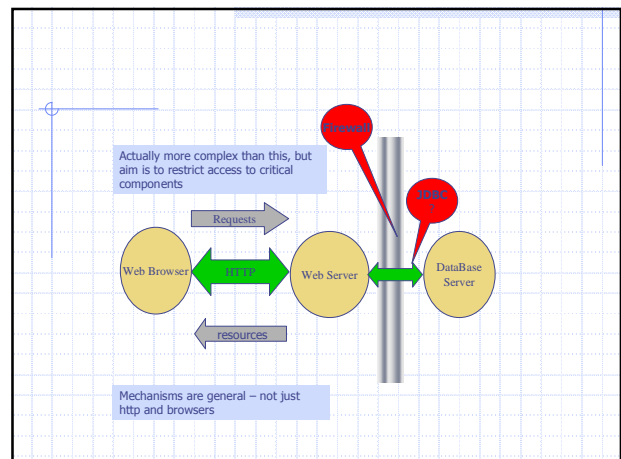
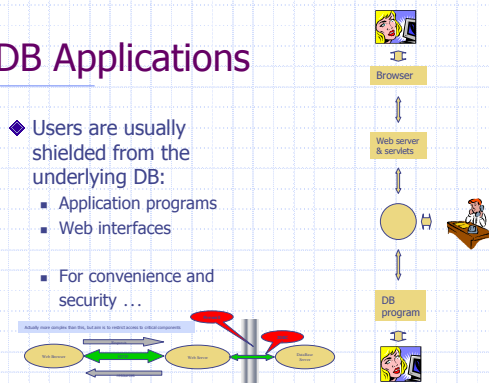
ANSI/SPARC DB Architecture

- ◆ View or external data model
 - A view onto (a subset of) conceptual model
- ◆ Conceptual data model
 - An abstract model independent of implementation
- ◆ Physical (internal) data model



DB Applications

- ◆ Users are usually shielded from the underlying DB:
 - Application programs
 - Web interfaces
- For convenience and security ...



Relational Databases

- ◆ Main model for data base systems
- ◆ Data in form of sets and relations
- ◆ Data connected using basic set theory
 - Selecting, combining etc.
- ◆ Normally *viewed as tables*
- ◆ Queries *specify* result **not** how it is computed
 - *declarative*

Example – just for illustration!

