

# Gradient Descent Learning (II): Classification

Shan He

School for Computational Science  
University of Birmingham

Module 06-27818 and 27819:  
Introduction to Neural Computation (Level 4/M)  
Neural Computation (Level 3)


# Outline of Topics

Classification

Logistic function

Solving Classification problems using ANNs

## An motivating example: Iris classification



## Classification: The Iris dataset

- ▶ The data set consists of 50 samples (training data) from each of three species of Iris:
  - ▶ Iris-setosa
  - ▶ Iris-virginica
  - ▶ Iris-versicolor
- ▶ Measures of four features of each samples in cm:
  - ▶ Petal length
  - ▶ Petal width
  - ▶ Sepal length
  - ▶ Sepal width
- ▶ Task: to train a model to learn from the training data to classify unseen Iris flowers based on the four features
- ▶ [Download the dataset](#) and [Detailed explanation of the dataset](#)


## Classification: a less interesting simple problem

- ▶ The Time spent and Math score/grade problem:

Student ID	1	2	3	4	5	6	7	8	9	10
Hours spent	4	9	10	14	4	7	12	22	3	10
Math Score	39	58	65	73	41	50	60	79	40	64
Grades	F	F	P	P	F	F	P	P	F	P

- ▶ Task: to train a classification model based on the hours spent and grades to predict unseen students' grades based on the hours spent

# Logistic function



# Classification Problems

- ▶ **Classification:** determining the most likely class that an input pattern belongs to.
- ▶ **Formally:** modelling the posterior probabilities of class membership conditioned on the input variables.
- ▶ Artificial neural networks: one output unit for each class, and for each input pattern we have
  - ▶ 1 for the output unit corresponding to that class
  - ▶ 0 for all the other output units
- ▶ The simplest case: binary classification  $\rightarrow$  one output unit
  - ▶ 1: passed
  - ▶ 0: failed

## Logistic regression

- ▶ **Logistic regression:** also known as logit regression, is a regression model where the prediction (dependent variable) is categorical, e.g., binary.
- ▶ **Goal:** to predict the probability that a given example belongs to the 1 class versus the probability that it belongs to the 0 class.



## Logistic regression

- **Logistic regression:** to learn a function of the form:


$$P(y = 1|\mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x})} \equiv \sigma(\mathbf{w}^T \mathbf{x})$$

and

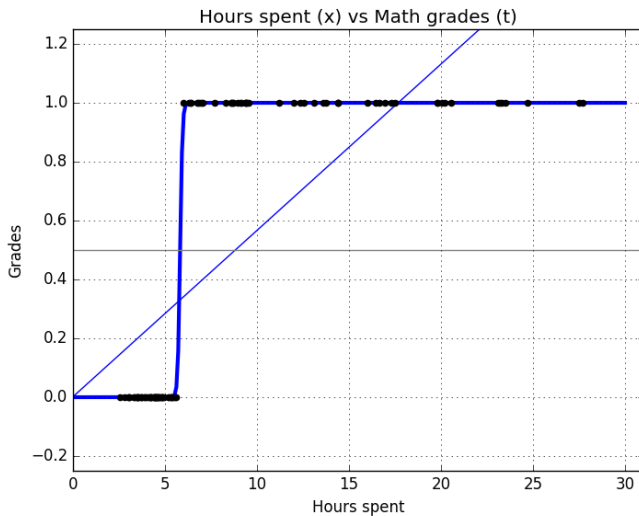
$$P(y = 0|\mathbf{x}) = 1 - P(y = 1|\mathbf{x})$$

- The function  $\sigma(x) = \frac{1}{1 + \exp(-x)}$  is often called the “logistic” or “sigmoid” function
- We usually use Maximum Likelihood estimation to obtain  $\mathbf{w}$

# Logistic function



# Logistic function




## Solving Classification problems using ANNs?

- ▶ How to solve this simple linear regression using neural networks.
- ▶ We need three things to define a neural networks:
  - ▶ **Network topology:** to define how neurons are connected by weights
  - ▶ **Activation function:** to convert a neuron's weighted input to its output activation
  - ▶ **Learning process:** to update the weights
    - ▶ **Essence of supervised learning process:** adjusting the network weights  $w_{ij}$  to minimise a **cost function**

## Solving classification problems using a perceptron?

- ▶ Network topology: Single Layer Networks
- ▶ Assuming binary classification: one output neuron



- ▶ Activation function
  - ▶ Question: can we use the simple linear activation function  $f(x) = x$ ?

## Solving classification problems using a perceptron?

- ▶ Activation function

- ▶ Question: can we use the simple linear activation function  $f(x) = x$ ?
  - ▶ Answer: No. Linear function outputs a continuous value, not a categorical output i.e., Iris-setosa or Iris-virginica, or Passed/Failed.
- ▶ We need different activation functions output categorical values:
  - ▶ Logistic function or Sigmoid function:

$$y = f(\mathbf{w}^T \mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x})} \equiv \sigma(\mathbf{w}^T \mathbf{x})$$

## Learning process

- ▶ **Learning process:** to update the weights
  - ▶ **Essence of supervised learning process:** adjusting the network weights  $w_{ij}$  to minimise a **cost function**

## Gradient Descent Learning

- ▶ **Aim:** to develop a learning algorithm that minimises a cost function (such as Sum Squared Error) by making appropriate **iterative adjustments** to the weights  $w_{ij}$ .
- ▶ **Idea:** to apply a series of small updates to the weights  $w_{ij}^{t+1} = w_{ij}^t + \Delta w_{ij}$  until the cost  $E(w_{ij})$  is “small enough”.
- ▶ **Question:** how to obtain  $\Delta w_{ij}$
- ▶ **Answer:**  $\Delta w_{ij} = -\eta \frac{\partial E(w_{ij})}{\partial w_{ij}}$
- ▶ **Explanation:** we repeatedly adjust the weights by small steps against the gradient, we will move through weight space, descending along the gradients towards a minimum of the cost function.



## Cross Entropy Cost Function for Two Classes

- ▶ Reminder: the Sum Squared Error cost function is for regression problems → we need a new cost function for classification problems
- ▶ If the output  $\mathbf{y}$  of a network represents the probability of a particular class, and  $\mathbf{t}$  is the **binary** target output, the probability of observing the whole training data set is:

$$P(\mathbf{t}|\mathbf{x}) = \prod_{p=1}^P y_p^{t_p} \cdot (1 - y_p)^{1-t_p}$$

- ▶ The ANN model aims to maximise this probability
- ▶ It is more convenient to work with a negative log-likelihood function since it is monotonically decreasing:

$$-\log \mathcal{L}(\mathbf{w}|\mathbf{t}, \mathbf{y}) = - \sum_p [t^p \log(y^p) + (1 - t^p) \log(1 - y^p)]$$

## Cross Entropy Cost Function for Two Classes

- ▶ We define the Cross Entropy Cost Function

$$\begin{aligned} E_{ce} &= -\log \mathcal{L}(\mathbf{w}|t, y) = -\sum_p [t^p \log(y^p) + (1 - t^p) \log(1 - y^p)] \\ &= -\sum_p [t^p \log(\sigma(\mathbf{w}^T \mathbf{x})) + (1 - t^p) \log(1 - \sigma(\mathbf{w}^T \mathbf{x}))] \end{aligned}$$

- ▶ The cost function looks very complicated but we shall see some tricks to derive its derivative

## Gradient Descent Learning for Cross Entropy Cost Function

- ▶ Cross Entropy Cost Function for the  $p$ -th training sample  $(\mathbf{x}^p, t^p)$ :

$$\Delta \mathbf{w}^p = -\eta \frac{\partial E_{ce}^p(\mathbf{w})}{\partial \mathbf{w}}$$

- ▶ Reminder: Using sigmoid activation:  $y^p = \sigma(\mathbf{w}^T \mathbf{x}^p)$ .
- ▶ Using the chain rule:

$$\begin{aligned}\Delta \mathbf{w}^p &= -\eta \frac{\partial E_{ce}^p(\mathbf{w})}{\partial \mathbf{w}} = -\eta \frac{\partial E_{ce}^p(\mathbf{w})}{\partial y^p} \frac{\partial y^p}{\partial \mathbf{w}} = -\eta \frac{\partial E_{ce}^p(\mathbf{w})}{\partial y^p} \frac{\partial \sigma(\mathbf{w}^T \mathbf{x}^p)}{\partial \mathbf{w}} \\ &= -\eta \frac{\partial E_{ce}^p(\mathbf{w})}{\partial y^p} \frac{\partial \sigma(\mathbf{w}^T \mathbf{x}^p)}{\partial \mathbf{w}^T \mathbf{x}^p} \frac{\partial \mathbf{w}^T \mathbf{x}^p}{\partial \mathbf{w}} = -\eta \frac{\partial E_{ce}^p(\mathbf{w})}{\partial y^p} \sigma' \mathbf{x}^p\end{aligned}$$

## The Derivative of a Sigmoid: $\sigma'$

- ▶ Let  $\sigma(x) = \frac{1}{1+\exp(-x)} = g(h(x))$  with  $g(h) = h^{-1}$  and  $h(x) = 1 + \exp(-x)$ , so

$$\frac{\partial g(h)}{\partial h} = -\frac{1}{h^2}, \quad \frac{\partial h(x)}{\partial x} = -\exp(-x)$$

- ▶ Using the chain rule:

$$\begin{aligned}\sigma' &= \frac{\partial \sigma(x)}{\partial x} = \frac{\partial g(h)}{\partial h} \frac{\partial h(x)}{\partial x} = -\frac{1}{h^2} \cdot -\exp(-x) = \frac{\exp(-x)}{(1 + \exp(-x))^2} \\ &= \frac{1 + \exp(-x) - 1}{(1 + \exp(-x))^2} = \frac{1}{(1 + \exp(-x))} - \frac{1}{(1 + \exp(-x))^2} \\ &= \sigma(x)(1 - \sigma(x))\end{aligned}$$

## Question: Why use Sigmoid function?

- ▶ Question: Why not use step function?

## Cross Entropy Cost Function for Two Classes

- We then derive

$$\begin{aligned}\frac{\partial E_{ce}^p}{\partial y^p} &= - \frac{\partial(t^p \log(y^p) + (1 - t^p) \log(1 - y^p))}{\partial y^p} \\ &= - \left[ \frac{\partial(t^p \log(y^p))}{\partial y^p} + \frac{\partial((1 - t^p) \log(1 - y^p))}{\partial y^p} \right]\end{aligned}$$

- Since:

$$\frac{\partial(t^p \log(y^p))}{\partial y^p} = \frac{\partial(t^p \log(y^p))}{\partial \log(y^p)} \frac{\partial \log(y^p)}{\partial y^p} = \frac{t^p}{y^p}$$

and

$$\frac{\partial((1 - t^p) \log(1 - y^p))}{\partial y^p} = - \frac{1 - t^p}{1 - y^p}$$

- Finally:

$$\frac{\partial E_{ce}^p}{\partial y^p} = - \frac{t^p}{y^p} + \frac{1 - t^p}{1 - y^p} = \frac{y^p - t^p}{y^p(1 - y^p)}$$

## Cross Entropy Cost Function for Two Classes

- ▶ Since  $y^p = \sigma(\mathbf{w}^T \mathbf{x}^p)$ , and we know:

$$\sigma'(\mathbf{w}^T \mathbf{x}^p) = \sigma(\mathbf{w}^T \mathbf{x}^p)(1 - \sigma(\mathbf{w}^T \mathbf{x}^p)) = y^p(1 - y^p)$$

- ▶ Putting them together:

$$\begin{aligned}\Delta \mathbf{w}^p &= -\eta \frac{\partial E_{ce}^p(\mathbf{w})}{\partial y^p} \sigma' \mathbf{x}^p = -\eta \frac{y^p - t^p}{y^p(1 - y^p)} y^p(1 - y^p) \mathbf{x}^p \\ &= \eta(t^p - y^p) \mathbf{x}^p\end{aligned}$$

- ▶ In the batch processing, we just add up all the gradients for each sample:

$$\Delta \mathbf{w} = \eta \sum_{p=1}^P (t^p - y^p) \mathbf{x}^p$$

## Conclusion

- ▶ We learned how to solve classification problems using a perceptron with sigmoid output activations and a Cross Entropy cost function
  - ▶ Bishop: Sections 3.1, 6.1, 6.7, 6.8, 6.9
  - ▶ Haykin-1999: Sections 3.5, 3.7, 4.4, 4.6