

Co-Evolution


Ata Kaban

University of Birmingham

Cheetah: 60-70 mph, for up to 100 yards



Thompson's Gazelle: 50 mph...



The outcome (but about 50% of attempts are failures)

What is Co-Evolution?

- Fitness of an individual depends on other individuals
 - Fitness landscape changes
 - Fitness of an individual may be different in different runs
- Change in one individual will change the fitness landscape in others
- Remember the Iterated Prisoner's Dilemma experiments? – the ones in which evolving strategies played against each other

Types of co-evolution

- By evaluation
 - Competitive
 - Cooperative
- By population-organisation
 - Inter-population
 - Intra-population

Example1: Sorting algorithm

- Goal: place the elements in a data structure (e.g. list or tree) in some specified order (e.g. numerical or alphanumeric)
- One approach in Knuth's book is the so called sorting network (for fixed nos of elements)
 - Horizontal lines=elements in the list
 - Vertical arrows=comparisons to be made (in parallel)
 - If compared elements are in wrong order than swap
- Will look at the simple case of $n=16$ elements (e0—e15)

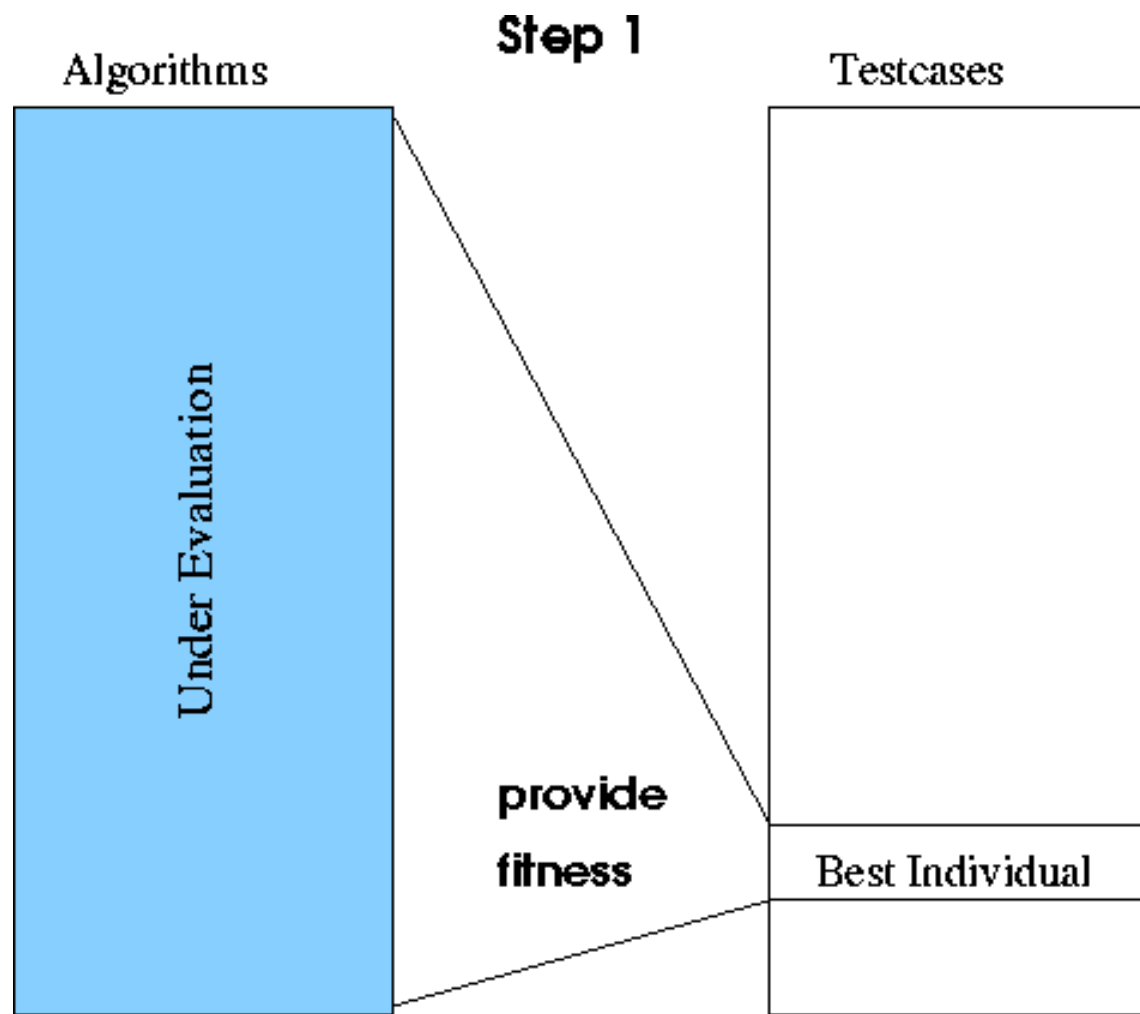
Designing sorting algorithms

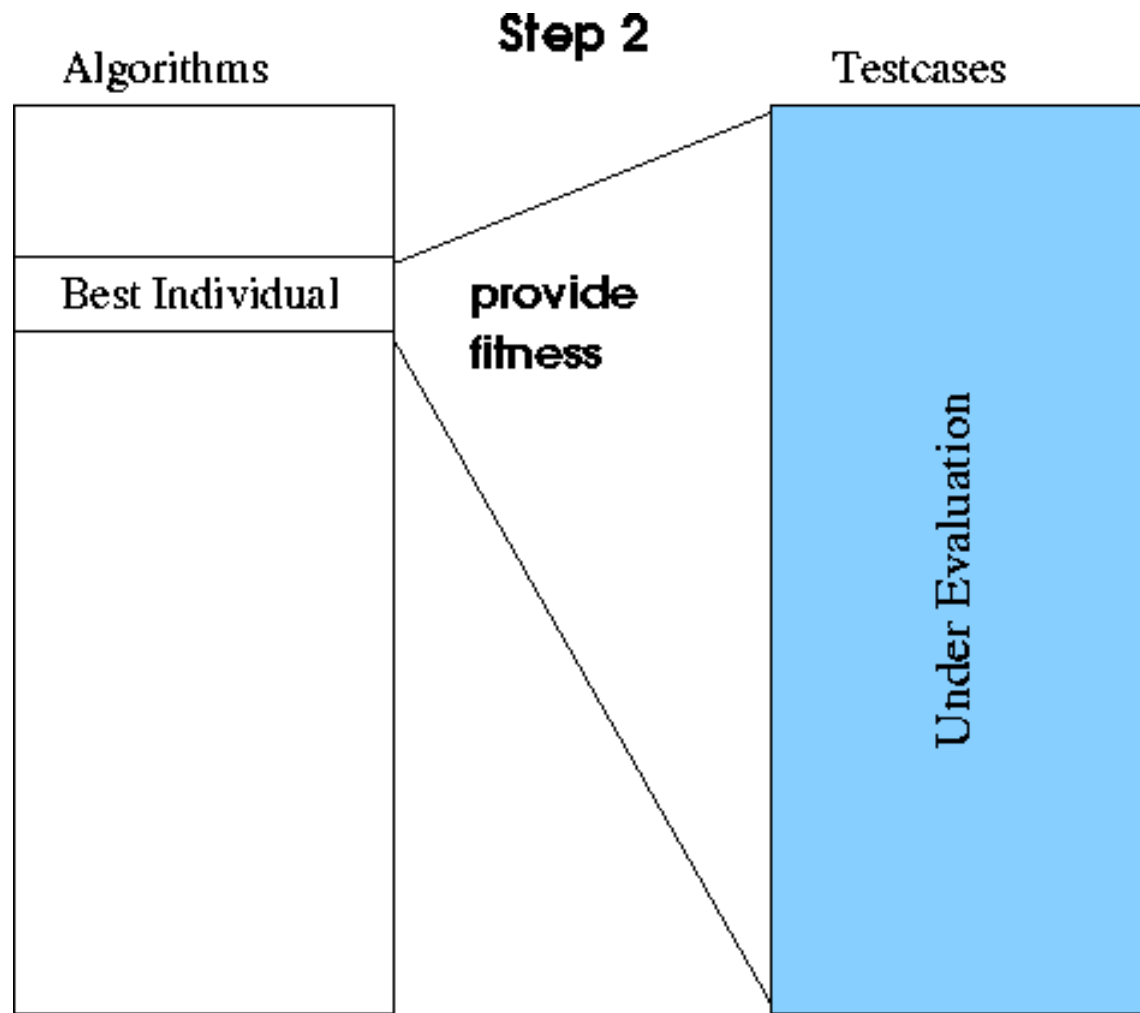
- Goal: make them
 - correct
 - efficient (reduce no of comparisons)
- What is the minimum number of comparisons necessary for correct sorting?
 - hot Q around the '60s ($n=16$ case)
 - '62: Bose & Nelson developed a net that needs 65 comparisons
 - '64: Batcher, Flyod, Knuth – 63 comparisons
 - '69: Shapiro – 62 comparisons
 - '69: Green – 60 comparisons
 - '80s: W.D. Hillis – can GA find an answer to this problem?

- Encoding the sorting net
 - ordered list of pairs to be compared = phenotypes
 - considered nets of 60—120 comparisons
 - genotype: diploid chromosomes
 - 1 individual = 15 pairs of 32 bit chromosomes, each encoding 4 comparisons
 - See more details on M. Mitchell, pp. 21—27.
- Fitness measure: the percentage of cases sorted correctly
 - Problem: how to compute this?
 - Test all possible inputs – slow
 - Test fixed set of inputs? – which?

- First trials
 - To foster speciation, individuals placed on a 2D grid (i.e. spatial distance between them)
 - Fitness computed from random subsamples
 - Half of population with lower fitness deleted – replaced with a copy of a surviving neighbor
 - Pairing in the local neighborhoods
 - Special crossover for diploids, followed by mutation with $p_m = 0.001$.
 - Population size between 512—1million
 - 5000 generation
- result: GA found sorting net of 65 comparisons
- ☹

- Why didn't the GA do better?
 - After early generations, with randomly generated test cases used to compute fitness, the difficulty of test cases stayed roughly the same!
- Solution: Co-Evolution
 - = evolve both algorithms and test cases!
 - Algorithms try to sort
 - Test cases try to 'trip up' algorithms
 - predator/prey (or host/parasite) relationship inspired from nature
 - as the algorithms got better, the test cases got harder (specifically targeting weaknesses in the networks)
- Result: 61 comparisons





- Importance of the work:
 - introduces new technique inspired by co-evolution. The results are convincing in that this new technique is potentially powerful.
- Inter-population competitive co-evolution

Example2: Game Playing

- Intra-population competitive co-evolution
- Task: evolve a backgammon player
- Problem: evaluation
 - Against human player
 - Against ‘conventional’ program
 - Against internet players
- Solution: co-evolution
 - Play against other evolving programs
- Intra-population
 - All genotypes are of the same type
 - Only one population

Co-Evolving Backgammon Players

- TD-Gammon
 - Grand Master level player
 - Learns by self-playing
- What makes it successful ?
 - Temporal-Difference Learning ?
 - Or simply self-playing ?
- Simple use of NN in game play:
 - Generate all legal moves
 - Feed them to the NN
 - Best output is new move
 - NN used to evaluate positions

Simple Backgammon Learner

Evolve a NN that plays backgammon

1. Initial NN is $NN(k)$; $k=0$;

2. Generate a mutant challenger of NN_k :

$$w'(i,j) = w(i,j) + \text{Gaussian}(0, s);$$

3. If $NN'(k)$ is beaten by $NN(k)$

$$NN(k+1) = NN(k)$$

else

$$NN(k+1) = 0.95 * NN(k) + 0.05 * NN'(k)$$



4. $k = k+1$; go to step 2 unless finished

- Setting
 - Uses 197-20-1 fully connected feed-forward NN; initial weights 0s
 - No training for NNs !
 - EA with population size 1
 - Only simple mutation: add Gaussian noise to weights
 - No recombination
 - Performance
- 40% winning against a strong expert-trained program (PUBEVAl) after 100.000 generations

Example3: Evolution of a complete robot

- Inter-population Cooperative Co-Evolution
- **Task: Evolve Morphology and Behaviour of a robot**
 - **Structure**
 - **Control**
- **Evaluation: no problem**
 - **Simulation**
 - **But: slow**
- **Problem: Evolution**
 - **All-in-one is possible, but:**
 - **Very large search space**
- **Solution: Co-Evolution**
 - **Divide Problems**
 - **Evolve Structure and Control separately**

GOLEM project, Pollack et. al



Coupling cooperative Co-Evolution

- **Tight coupling (1)**
 - Body-Individual and Controller-Individuals evolve into 'matched pair'
 - Body 'a' only works well with controller 'b'
- **Option: Two-Part Genotype**
 - Keep both genotypes in one Individual
 - Evolve both together
- **Tight coupling (2)**
 - Change in one might depend on change in the other
 - Can slow down evolution
- **Option: delayed evaluation**
 - Mutate morphology individuals
 - Evolve controller individuals (possibly more than one generation)
 - Assign fitness to morphology individuals

Example4: Pattern recognition

- **Intra-population Cooperative Co-Evolution**
- **Taks: Evolve a NN to recognize a set of letters**
 - Grey values fed into NN
 - Output classifies letters
 - Classification problem
- **Co-Evolution: Divide and Conquer**
 - Different Individuals specialize on different letters
 - Group output is combined into one by external mechanism (e.g. voting)
- **Speciation**

When and Why Co-Evolution?

- **No fitness-function known**
 - **Bootstrapping by co-evolution**
- **Too many fitness cases**
 - **Co-Evolve fitness and cases**
- **Modularizable Problem**
 - **Divide and Conquer**

Other examples

- **Creative Design Systems**
 - Evolve designs and design specifications
- **Other Games**
 - Tic-Tac-Toe
 - Prisoner's Dilemma
 - Checkers
- **Artificial Life**
 - Complex simulated Ecosystems

Summary

- Co-Evolution is Everywhere
- Can be cooperative and competitive
- Can be in one population, or more than one
- An individual's fitness is not fixed in co-evolution
- Co-evolution is not well explored in Evolutionary Computation

References

- Mitchell, M, Introduction to Genetic Algorithms (book), MIT Press, 1996, pp.21—27.
- Pollack, J, Blair, A. and Land, M. (1996). Coevolution of A Backgammon Player. *Proceedings Artificial Life V*, C. Langton, (Ed), MIT Press. Available at: <http://www.demo.cs.brandeis.edu/papers/alife5.pdf>
- Pollack, Jordan B., Lipson, Hod, Hornby, Gregory S., and Funes, Pablo. Three Generations of Automatically Designed Robots. *Artificial Life*, 7:3, pg 215-223. 2001. Available at: http://www.demo.cs.brandeis.edu/papers/long.html#pollack_alife01
- Darwen, P. J. and Yao, X. (1996). Automatic Modularization by Speciation. *Third IEEE International Conference on Evolutionary Computation*, available at: <http://www.demo.cs.brandeis.edu/papers/icec96darwen.ps.gz>