

Stochastic Local Search Algorithms for TSP

Shan He

School for Computational Science
University of Birmingham

Module 06-27818 and 27819: Advanced Aspects of
Nature-Inspired Search and Optimisation (Ext)

Outline of Topics

- 1 Local search algorithms
 - Hill-climbing search
- 2 2-Opt algorithm
- 3 Exploration vs exploitation
- 4 Stochastic local search: basic ideas
- 5 Simulated Annealing

Question

Why our randomise algorithm didn't work on TSP?

Problems with randomised search algorithms

- Q: Why randomised search doesn't work on TSP??
- A: In TSP or any other optimisation problems, acceptable near optimal solutions are just a small proportion of all possible solutions
- Q: How to design algorithms to deal with this kind of problems?
- A: Most of the optimisation problems, including TSP, have a highly coherent search (solution) space, i.e., the solutions have **neighbourhood structure**.
- **Neighbourhood structure**: each solution has at least one neighbourhood solution, i.e., similar solutions
- We can exploit this coherence to design heuristic algorithms

Problems with randomised search algorithms

- A visualisation of the search landscape of a difficult combinatorial optimisation problem: **social network community detection problems**

Local search algorithms

- **Local search:** a metaheuristic method for solving hard optimization problems
 - **Idea:** start with an initial guess at a solution and **incrementally** improve it until it is one
 - **Incremental improvement:** local changes, e.g., the algorithm iteratively moves to a **neighbour solution**
 - **Neighbour solution:** Depends on the definition of a neighbourhood relation on the search space, but generally based on similarity (distance) measure

Generic local search algorithm

Generic local search algorithm

```
 $x_0$  := generate initial solution
terminationflag := false
 $x$  :=  $x_0$ 
while (terminationflag != true)
    Modify the current solution to a neighbour one  $v \in \mathcal{A}$ 
    If  $f(v) < f(x)$  then  $x := v$ 
    If a termination criterion is met: terminationflag := true
Output  $x$ 
```

Note: termination criterion could be maximum iteration is reached or no improvement for a certain iterations.

Hill climbing algorithm

- One of the simplest local search algorithms
- Hill climbing is an algorithm that more like “climbing Everest in thick fog with amnesia”
- An iterative algorithm:
 - Starts with an arbitrary solution to a problem,
 - Iteratively searches a better solution from the current solution's **immediate neighbour solutions**
 - **Immediate neighbour solutions**: most similar solutions to the current solution.
- Two types of hill climbing:
 - **Simple hill climbing**: chooses the **first** better solution
 - **Steepest ascent hill climbing**: compares all neighbour solutions and chooses **the best** solution

Simple hill climbing algorithm

Simple hill climbing algorithm

x_0 := generate initial solution

terminationflag := false

x := x_0

while (terminationflag != true)

 Modify the current solution to a **immediate** neighbour one $v \in \mathcal{A}$

 If $f(v) < f(x)$ then $x := v$

 If a termination criterion is met: terminationflag := true


Output x

Hill climbing for TSP problem

- Question: How to construct the **immediate neighbour solutions** of the current solution for TSP?

Let's take a look at some simple examples

- 2-3 cities: only one solutions
- 4 cities: 3 solutions
- **Question:** How those tours of the 4 cities TSP differ?



2-Opt algorithm


- 2-Opt: A simple local search operator first proposed by Croes in 1958 [1] for solving the travelling salesman problem.
- Basic idea: **two immediate neighbour solutions can be two routes (cycles) only differ from two edges**
- Swapping two edges can result in two **immediate** neighbour solutions
- Detailed swapping steps:
 - Step 1: removal of two edges from the current route, which results in two parts of the route.
 - Step 2: reconnect by two other edges to obtain a new solution

[1] G.A. Croes. A method for solving traveling-salesman problems. Operations Research. 1958

2-Opt algorithm


Suppose we have a route: $A \rightarrow C \rightarrow B \rightarrow D \rightarrow A$, which is obviously not optimal. Let's see how to swap:

Step 1: removal of two edges from the current route, which results in two parts of the route



2-Opt algorithm

Step 2: reconnect by two other edges to obtain a new solution.



This will result in an optimal route: $A \rightarrow B \rightarrow C \rightarrow D \rightarrow A$



How to implement 2-Opt algorithm?

- Compare the two solutions:
 - $A \rightarrow C \rightarrow B \rightarrow D \rightarrow A$
 - $A \rightarrow B \rightarrow C \rightarrow D \rightarrow A$
- We observed that we swap two **adjacent** cities, e.g., B and C in a route can create two **immediate** neighbour solutions
- But how about this 6 cities TSP problem?

2-Opt algorithm



Suppose we have a route: $A \rightarrow B \rightarrow C \rightarrow D \rightarrow E \rightarrow F \rightarrow A$, which is obviously not optimal. Let's see how to swap:

Step 1: removal of two edges from the current route, which results in two parts of the tour



2-Opt algorithm

Step 2: reconnect by two other edges to obtain a new solution



We need to reverse the order of $B \leftarrow C \leftarrow D \leftarrow E$ in order to get
 $A \rightarrow E \rightarrow D \rightarrow C \rightarrow B \rightarrow F \rightarrow A$

2-Opt algorithm: implementation

2-Opt algorithm

route := initial TSP solution

i, j := two cities for swapping

Step 1: take *route*[1] to *route*[*i* − 1] and add them in order to *newroute*

Step 2: take *route*[*i*] to *route*[*k*] and add them in reverse order to *newroute*

Step 3: take *route*[*k* + 1] to end and add them in order to new *newroute*

Output *newroute*

Exercise: Simple hill climbing for TSP

- Implement the 2-Opt algorithm by editing my empty twoopt.m file

- First test your implementation using

```
test_city = ['A', 'B', 'C', 'D', 'E', 'F'];  
newroute = twoopt(test_city, 2, 5)
```

- Once you validate your 2-opt algorithm, then execute:

```
clear all  
load('cities.mat')  
[dist route] = simple_hill_climbing_two_opt(cities)
```

- Execute

```
[dist route] = simple_hill_climbing_two_opt(cities)
```

a few times, what is your observation?


- Modify the simple hill climbing algorithm as the steepest ascent/descent hill climbing algorithm.

Issues about hill climbing

- We can observe: the algorithm got stuck at an improved solution (local optimum) but still not the best (global optimum)
- Global vs local optimum
 - **Local optimum:** a solution that is optimal (either maximal or minimal) within a neighbouring set of candidate solutions
 - **Global optimum:** the optimal solution among all possible solutions
- General issues with local search:
 - Find only local optima unless the search space is unimodal (only one optima)
 - Even for unimodal problem, if there exists plateau, hill climbing might not perform well.

Fitness landscape with global/local optima

Fitness landscape of a 1-dimensional optimisation problem



Intensification vs Diversification

- **Intensification** (Exploitation): search small region around the current solution
 - Aim to improve a promising solution S that we already have at hand by searching in the vicinity of S
 - Local search \longrightarrow local optimum
- **Diversification** (Exploration): search large unknown region of the search space
 - Aim to find other promising solutions that are yet to be refined
 - Need to escape from current local optimum \longrightarrow Randomness can help
 - Global search \longrightarrow global optimum

Randomised search vs Local search

- **Randomised search:**

- Good at exploration, e.g., to search large unknown region of the search space
- Not good at exploitation, e.g., to search small region around the current solution
- Especially bad for problems where good solutions are just a small portion of all possible solutions

- **Local search:**


- Good at exploitation: capable to find local optimum
- Not good at exploration: gets stuck into local optimum
- Question: can we combine these two heuristics?

Stochastic local search: Main idea

- Main idea: escape or avoid (more difficult) local optima
- How: Introduce randomness into local search algorithm to escape from local optima
- Escape strategies:
 - **Random restart**: simply restart the local search from a random initial solution
 - Applicable when:
 1. Number of local optima is small
 2. The cost for restarting the local search is cheap
 - **Perform random non-improving step**: randomly move to a less fit neighbour

Search trajectory of stochastic local search

Minimisation problem, e.g., TSP problem.



Stochastic Simple hill climbing algorithm

Stochastic simple hill climbing algorithm

$x_0 :=$ generate initial solution

terminationflag := false

$x := x_0$

while (terminationflag != true)

 Modify the current solution to a **immediate** neighbour one $v \in \mathcal{A}$

 If $f(v) < f(x)$ or $\text{rand}(1) < P$ then $x := v$

 If a termination criterion is met: terminationflag := true

Output x

Note: P is a small probability.

Exercises: Stochastic simple hill climbing for TSP

- Modify the Simple Hill Climbing Search Algorithm code to construct a Stochastic Hill Climbing Search Algorithm.
 - Hint 1: instead of rejecting worse solutions (immediate neighbours), accept them with a very small probability, e.g., 0.001.
 - You also need to specify the maximum iteration, otherwise your algorithm will never terminate

Simulated Annealing (Minimisation, e.g., TSP problem)

- Main idea: Escape from local optima with random non-improving step:
 - Accepting good solution with the probability of 1, e.g., $P := 1$ if $f(v) < f(x_t)$
 - Accepting worse solution with a certain probability, e.g.,
$$P := e^{-\frac{f(v) - f(x_t)}{T(t)}} \text{ if } f(v) \geq f(x_t)$$
- The annealing schedule $T(t)$ will slowly decrease in the probability of accepting worse solutions
- You can implement your SA from the Stochastic simple hill climbing algorithm

Generic Simulated Annealing algorithm for minimisation

```

 $x := x_0; e := f(x)$       // Initial solution, objective function value (energy).
 $x_{best} := x; e_{best} := e$   // Initial "best" solution
 $k := 0$       // Count evaluation number.
while ( $k < k_{max}$ )
     $T := temperature(t_0)$     // Temperature calculation.
     $x_{new} := neighbour(x)$     // Pick some neighbour.
     $e_{new} := f(x_{new})$       // Compute its objective function value.
    if  $P(e, e_{new}, T) > R(0, 1)$  then    // Should we move to it?
         $x := x_{new}; e := e_{new}$     // Yes, change state.
    if  $e_{new} < e_{best}$  then    // Is this a new best?
         $x_{best} := x_{new}; e_{best} := e_{new}$     // Save as 'best found'.
     $k := k + 1$  // Increase Evaluation
Output  $x_{best}$ 

```

$P := 1$ if $e_{new} < e$, and

$P := e^{\frac{e - e_{new}}{T}}$ otherwise

Annealing schedule $temperature()$ defines how to decreased temperature from a initial temperature t_0 .